

KPMG TASK-1

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

Exploring Transactions Sheet

```
In [2]: path=r'C:\Users\JAYADEVA JAVALI\Downloads\KPMG_VI_New_raw_data_update_final.xlsx'
df=pd.read_excel(path,sheet_name='Transactions')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	pro
0	1	2	2950	2017-02-25	0.0	Approved	Solex	
1	2	3	3120	2017-05-21	1.0	Approved	Trek Bicycles	
2	3	37	402	2017-10-16	0.0	Approved	OHM Cycles	
3	4	88	3135	2017-08-31	0.0	Approved	Norco Bicycles	
4	5	78	787	2017-10-01	1.0	Approved	Giant Bicycles	

```
In [4]: df.shape
```

```
Out[4]: (20000, 13)
```

```
In [5]: df.columns
```

```
Out[5]: Index(['transaction_id', 'product_id', 'customer_id', 'transaction_date',
               'online_order', 'order_status', 'brand', 'product_line',
               'product_class', 'product_size', 'list_price', 'standard_cost',
               'product_first_sold_date'],
              dtype='object')
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   transaction_id                        20000 non-null  int64
1   product_id                          20000 non-null  int64
2   customer_id                         20000 non-null  int64
3   transaction_date                    20000 non-null  datetime64[ns]
4   online_order                        19640 non-null  float64
5   order_status                        20000 non-null  object
6   brand                              19803 non-null  object
7   product_line                       19803 non-null  object
8   product_class                      19803 non-null  object
9   product_size                       19803 non-null  object
10  list_price                          20000 non-null  float64
```

```

11 standard_cost          19803 non-null float64
12 product_first_sold_date 19803 non-null float64
dtypes: datetime64[ns](1), float64(4), int64(3), object(5)
memory usage: 2.0+ MB

```

```

In [10]: #convert date columns from integer to datetime
df['transaction_date'] = pd.to_datetime(df['transaction_date'], unit='s')
df['transaction_date'].head()

```

```

Out[10]: 0    2017-02-25
         1    2017-05-21
         2    2017-10-16
         3    2017-08-31
         4    2017-10-01
Name: transaction_date, dtype: datetime64[ns]

```

```

In [11]: #convert date columns from integer to datetime
df['product_first_sold_date'] = pd.to_datetime(df['product_first_sold_date'], unit='s')
df['product_first_sold_date'].head()

```

```

Out[11]: 0    1970-01-01 11:27:25
         1    1970-01-01 11:35:01
         2    1970-01-01 10:06:01
         3    1970-01-01 10:02:25
         4    1970-01-01 11:43:46
Name: product_first_sold_date, dtype: datetime64[ns]

```

```

In [7]: df.describe()

```

```

Out[7]:
```

	transaction_id	product_id	customer_id	online_order	list_price	standard_cost	produc
count	20000.000000	20000.000000	20000.000000	19640.000000	20000.000000	19803.000000	
mean	10000.500000	45.36465	1738.246050	0.500458	1107.829449	556.046951	
std	5773.647028	30.75359	1011.951046	0.500013	582.825242	405.955660	
min	1.000000	0.000000	1.000000	0.000000	12.010000	7.210000	
25%	5000.750000	18.000000	857.750000	0.000000	575.270000	215.140000	
50%	10000.500000	44.000000	1736.000000	1.000000	1163.890000	507.580000	
75%	15000.250000	72.000000	2613.000000	1.000000	1635.300000	795.100000	
max	20000.000000	100.000000	5034.000000	1.000000	2091.470000	1759.850000	

```

In [9]: df.isnull().sum()

```

```

Out[9]: transaction_id          0
        product_id            0
        customer_id           0
        transaction_date       0
        online_order          360
        order_status          0
        brand                 197
        product_line          197
        product_class         197
        product_size          197
        list_price             0
        standard_cost         197
        product_first_sold_date 197
        dtype: int64

```

To Treat Missing Values, The options we have are:

- **Drop Missing Values**
- **Impute Missing Values based on type of variable**

We can decide on this during analysis based on objective.

```
In [12]: dups = df.duplicated()
         dups.sum()
```

```
Out[12]: 0
```

Exploring Columns

```
In [13]: df['order_status'].value_counts()
```

```
Out[13]: Approved      19821
         Cancelled      179
         Name: order_status, dtype: int64
```

```
In [14]: df['brand'].value_counts()
```

```
Out[14]: Solex          4253
         Giant Bicycles  3312
         WeareA2B        3295
         OHM Cycles      3043
         Trek Bicycles   2990
         Norco Bicycles  2910
         Name: brand, dtype: int64
```

```
In [15]: df['product_line'].value_counts()
```

```
Out[15]: Standard      14176
         Road          3970
         Touring        1234
         Mountain       423
         Name: product_line, dtype: int64
```

```
In [16]: df['product_class'].value_counts()
```

```
Out[16]: medium      13826
         high         3013
         low          2964
         Name: product_class, dtype: int64
```

```
In [17]: df['product_size'].value_counts()
```

```
Out[17]: medium      12990
         large        3976
         small        2837
         Name: product_size, dtype: int64
```

```
In [54]: df['transaction_id'].nunique()
```

```
Out[54]: 20000
```

```
In [58]: df['customer_id'].nunique()
```

```
Out[58]: 3494
```

```
In [18]: df.corr().T
```

```
Out[18]:
```

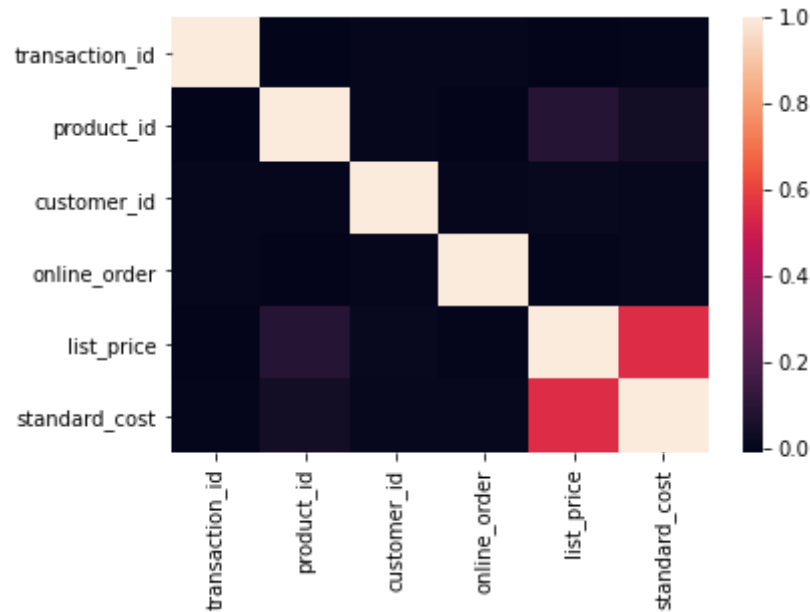
	transaction_id	product_id	customer_id	online_order	list_price	standard_cost
--	----------------	------------	-------------	--------------	------------	---------------

	transaction_id	product_id	customer_id	online_order	list_price	standard_cost
transaction_id	1.000000	-0.011486	0.001753	0.003394	-0.006154	-0.003291
product_id	-0.011486	1.000000	0.004278	-0.004233	0.090066	0.038765
customer_id	0.001753	0.004278	1.000000	0.001616	0.009306	0.005365
online_order	0.003394	-0.004233	0.001616	1.000000	-0.000295	0.006934
list_price	-0.006154	0.090066	0.009306	-0.000295	1.000000	0.551539
standard_cost	-0.003291	0.038765	0.005365	0.006934	0.551539	1.000000

In [20]:

sns.heatmap(df.corr())

Out[20]: <AxesSubplot:>



Exploring Customer Demographic

In [22]:

df2=pd.read_excel(path,sheet_name='CustomerDemographic')

In [23]:

df2.head()

Out[23]:

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job
0	1	Laraine	Medendorp	F	93	1953-10-12	Executive Secretary
1	2	Eli	Bockman	Male	81	1980-12-16	Administrative Clerk
2	3	Arlin	Dearle	Male	61	1954-01-20	Recreation Manager
3	4	Talbot	NaN	Male	33	1961-10-03	
4	5	Sheila-kathryn	Calton	Female	56	1977-05-13	Senior Manager

In [24]: `df2.shape`

Out[24]: (4000, 13)

In [25]: `df2.columns`

Out[25]: Index(['customer_id', 'first_name', 'last_name', 'gender', 'past_3_years_bike_related_purchases', 'DOB', 'job_title', 'job_industry_category', 'wealth_segment', 'deceased_indicator', 'default', 'owns_car', 'tenure'], dtype='object')

In [26]: `df2.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 13 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   customer_id                          4000 non-null   int64
 1   first_name                           4000 non-null   object
 2   last_name                            3875 non-null   object
 3   gender                               4000 non-null   object
 4   past_3_years_bike_related_purchases 4000 non-null   int64
 5   DOB                                  3913 non-null   datetime64[ns]
 6   job_title                            3494 non-null   object
 7   job_industry_category                3344 non-null   object
 8   wealth_segment                       4000 non-null   object
 9   deceased_indicator                   4000 non-null   object
10   default                              3698 non-null   object
11   owns_car                             4000 non-null   object
12   tenure                               3913 non-null   float64
dtypes: datetime64[ns](1), float64(1), int64(2), object(9)
memory usage: 406.4+ KB
```

In [27]: `df2.describe()`

	customer_id	past_3_years_bike_related_purchases	tenure
count	4000.000000	4000.000000	3913.000000
mean	2000.500000	48.890000	10.657041
std	1154.844867	28.715005	5.660146
min	1.000000	0.000000	1.000000
25%	1000.750000	24.000000	6.000000
50%	2000.500000	48.000000	11.000000
75%	3000.250000	73.000000	15.000000
max	4000.000000	99.000000	22.000000

In [55]: `df2['customer_id'].nunique()`

Out[55]: 4000

In [28]: `df2.isnull().sum()`

customer_id	0
first_name	0
last_name	125
gender	0
past_3_years_bike_related_purchases	0
DOB	87

```

job_title          506
job_industry_category  656
wealth_segment      0
deceased_indicator  0
default            302
owns_car            0
tenure             87
dtype: int64

```

To Treat Missing Values, The options we have are:

- **Drop Missing Values**
- **Impute Missing Values based on type of variable** We can decide on this during analysis based on objective.

```
In [31]: df2['gender'].value_counts()
```

```

Out[31]: Female    2037
Male      1872
U          88
Femal      1
F           1
M           1
Name: gender, dtype: int64

```

```
In [33]: # Replace inconsistent values with appropriate value
df2['gender'] = df2['gender'].replace('F', 'Female').replace('M', 'Male').replace('Fem', 'Female')
```

```
In [34]: df2['gender'].value_counts()
```

```

Out[34]: Female    2039
Male      1873
Unspecified    88
Name: gender, dtype: int64

```

```
In [35]: df2['default'].value_counts()
```

```

Out[35]: 100          113
1          112
-1         111
-100        99
â°â´âµâââ    53
...
ç¸¸æç¸¸â,éçèªâ.ç ç@¶æ    31
/dev/null; touch /tmp/blns.fail ; echo    30
âªâªtestâª    29
ì,ëë°í ë¥´    27
,ãã»:~:ã»ãã( â» Ì â» )ãã»:~:ã»ãã    25
Name: default, Length: 90, dtype: int64

```

The data doesnt seem to be right to process,so lets drop this column

```
In [37]: df2.drop('default',axis=1)
df2.head()
```

```

Out[37]:   customer_id  first_name  last_name  gender  past_3_years_bike_related_purchases  DOB  job
0          1    Laraine   Medendorp   Female                                93  1953-10-12  Executive
1          2         Eli    Bockman    Male                                81  1980-12-16  Administrator

```

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job
2	3	Arlin	Dearle	Male	61	1954-01-20	Recr Ma
3	4	Talbot	NaN	Male	33	1961-10-03	
4	5	Sheila-kathryn	Calton	Female	56	1977-05-13	Senior I

In [29]:


df2.corr()

Out[29]:

	customer_id	past_3_years_bike_related_purchases	tenure	
	customer_id	1.000000	-0.002529	-0.019947
	past_3_years_bike_related_purchases	-0.002529	1.000000	-0.009508
	tenure	-0.019947	-0.009508	1.000000

In [30]:

sns.heatmap(df.corr())



Exploring Customer Address

In [61]:

df3=pd.read_excel(path,sheet_name='CustomerAddress')
df3.head()

Out[61]:

	customer_id	address	postcode	state	country	property_valuation
0	1	060 Morning Avenue	2016	New South Wales	Australia	10
1	2	6 Meadow Vale Court	2153	New South Wales	Australia	10
2	4	0 Holy Cross Court	4211	QLD	Australia	9
3	5	17979 Del Mar Point	2448	New South Wales	Australia	4
4	6	9 Oakridge Court	3216	VIC	Australia	9

```
In [40]: df3.shape
```

```
Out[40]: (3999, 6)
```

```
In [41]: df3.columns
```

```
Out[41]: Index(['customer_id', 'address', 'postcode', 'state', 'country',
               'property_valuation'],
              dtype='object')
```

```
In [42]: df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 0 to 3998
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customer_id           3999 non-null   int64
1   address               3999 non-null   object
2   postcode              3999 non-null   int64
3   state                 3999 non-null   object
4   country               3999 non-null   object
5   property_valuation    3999 non-null   int64
dtypes: int64(3), object(3)
memory usage: 187.6+ KB
```

```
In [43]: df3.describe()
```

```
Out[43]:
```

	customer_id	postcode	property_valuation
count	3999.000000	3999.000000	3999.000000
mean	2003.987997	2985.755939	7.514379
std	1154.576912	844.878364	2.824663
min	1.000000	2000.000000	1.000000
25%	1004.500000	2200.000000	6.000000
50%	2004.000000	2768.000000	8.000000
75%	3003.500000	3750.000000	10.000000
max	4003.000000	4883.000000	12.000000

```
In [57]: df3['customer_id'].nunique()
```

```
Out[57]: 3999
```

```
In [53]: df3.isnull().sum()
```

```
Out[53]: customer_id      0
address      0
postcode     0
state        0
country      0
property_valuation  0
dtype: int64
```

There are no missing values here

```
In [48]: df3['address'].value_counts()
```

```
Out[48]: 3 Talisman Place      2
```



```
3 Mariners Cove Terrace      2
64 Macpherson Junction      2
205 Melody Circle            1
376 Buena Vista Street      1
..
590 Hayes Court              1
4365 Basil Junction          1
34748 Charing Cross Point    1
61 Kim Avenue                1
1 Cordelia Alley             1
Name: address, Length: 3996, dtype: int64
```

In [49]:

df3['postcode'].value_counts()

```
Out[49]: 2170      31
          2155      30
          2145      30
          2153      29
          2770      26
          ..
          4552       1
          4555       1
          2485       1
          3580       1
          4421       1
Name: postcode, Length: 873, dtype: int64
```

In [50]:

df3['state'].value_counts()

```
Out[50]: NSW          2054
          VIC          939
          QLD          838
          New South Wales    86
          Victoria         82
Name: state, dtype: int64
```

In [51]:

df3['country'].value_counts()

```
Out[51]: Australia    3999
Name: country, dtype: int64
```

In [52]:

df3['property_valuation'].value_counts()

```
Out[52]: 9      647
          8      646
          10     577
          7      493
          11     281
          6      238
          5      225
          4      214
          12     195
          3      186
          1      154
          2      143
Name: property_valuation, dtype: int64
```

In [44]:

df3.corr().T

Out[44]:

	customer_id	postcode	property_valuation
customer_id	1.000000	0.011396	-0.012073
postcode	0.011396	1.000000	-0.508392
property_valuation	-0.012073	-0.508392	1.000000

```
In [47]: sns.heatmap(df.corr())
```

```
Out[47]: <AxesSubplot:>
```

